

Home

Overview

Publications

Research

Software

▶ NAMD

▶ VMD

▶ GPU Computing

▶ Lattice Microbes

▶ Atomic Resolution
Brownian Dynamics

▶ MDFF

▶ QwikMD

▶ Other

Outreach

NAMD 3.0 Alpha, GPU-Resident Single-Node-Per-Replicate Test Builds

This page contains special developmental early "alpha" test builds of NAMD 3.0 intended to provide very early testing access to the NAMD user and developer communities.

- [NAMD 3 alpha test downloads](#)
- [NAMD 3 usage notes](#)

An exciting NAMD 3.0 feature being developed is a GPU-resident single-node-per-replicate computation mode that can speed up simulations up to two or more times for modern GPU architectures, e.g., Volta, Turing, Ampere.

This new NAMD simulation mode maximizes performance of small- to moderate-sized molecular dynamics simulations suitable for the computational capabilities of a single GPU-accelerated compute node, by virtue of new GPU-specific streamlined code paths that offload the integrator and rigid bond constraints to the GPU, while bypassing and thereby eliminating CPU activities and associated overheads that would otherwise slow down these simulations. With current developmental test builds of NAMD 3.0 alpha, the new code paths can provide a 2x performance gain compared to previous versions of NAMD running the same simulation on the same hardware.

What MD Simulations Are Well-Suited to NAMD 3.0 Alpha Versions?

This scheme is intended for small to medium systems (10 thousand to 1 million atoms). For larger simulations, you should stick to the regular integration scheme, e.g., as used in NAMD 2.x.

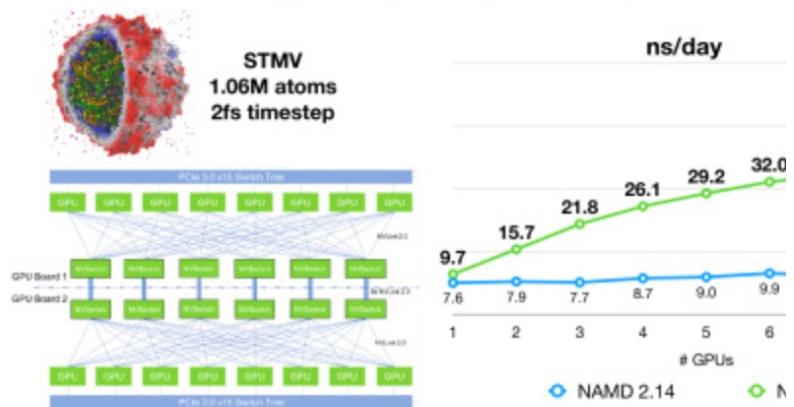
This scheme is intended for modern GPUs, and it might slow your simulation down if you are not running on a Volta, Turing, or Ampere GPU! If your GPU is older, we recommend that you stick to NAMD 2.x.

The single-node version of NAMD 3.0 has almost everything offloaded to the GPU, so large CPU core counts are NOT necessary to get good performance. We recommend running NAMD with a low +p count, maybe 2-4 depending on system size, especially if the user plans on running multiple replica simulations within a node.

NAMD 3.0 Feature Notes:

NAMD contains a lot of code! We currently don't have corresponding GPU code for everything that the original CPU is capable of handling. The consequence is that some features are still not supported, and others are supported but don't gain much performance (yet). Most of the supported features are related to equilibrium simulations, but some biasing schemes

NAMD 3.0: Single trajectory - Multiple GPU I



Early multi-GPU-per-replicate scaling plot for ST on an NVIDIA DGX-2 w/ Tesla V100 GPUs.

NAMD 3.0 Alpha: Performance for Apolipop



The following advanced features are not yet supported by multi-GPU runs: group position restraints, steered molecular dynamics (SMD), harmonic restraints, and electric field.

The code is still evolving, and test builds will be updated frequently. Stay tuned!

NAMD 3.0 Single- and Multi-GPU-Per-Replicate Alpha Versions For Download:

These are the first builds that support both single- and multi-GPU-per replicate runs in the same binary.

NAMD 3.0 alpha 9 and later support both single- and multi-GPU parallel scaling per replicate

- [NAMD_3.0alpha13_Linux-x86_64-multicore-CUDA-SingleNode.tar.gz](#) (Standard simulation.) (**July 24, 2022**)
- [NAMD_3.0alpha13_Linux-x86_64-netlrts-smp-CUDA-SingleNode.tar.gz](#) (Multi-copy GPU simulations. For multi-GPU machines use: "+devicesperreplica 1") (**July 24, 2022**)
- [NAMD_3.0alpha12_Linux-x86_64-multicore-CUDA-SingleNode.tar.gz](#) (Standard simulation.) (**June 22, 2022**)
- [NAMD_3.0alpha12_Linux-x86_64-netlrts-smp-CUDA-SingleNode.tar.gz](#) (Multi-copy GPU simulations. For multi-GPU machines use: "+devicesperreplica 1") (**June 22, 2022**)
- [NAMD_3.0alpha11_Linux-x86_64-multicore-CUDA-SingleNode.tar.gz](#) (Standard simulation.) (**May 24, 2022**)
- [NAMD_3.0alpha11_Linux-x86_64-netlrts-smp-CUDA-SingleNode.tar.gz](#) (Multi-copy GPU simulations. For multi-GPU machines use: "+devicesperreplica 1") (**May 24, 2022**)
- [NAMD_3.0alpha10_Linux-x86_64-multicore-CUDA-SingleNode.tar.gz](#) (Standard simulation.) (**March 24, 2022**)
- [NAMD_3.0alpha10_Linux-x86_64-netlrts-smp-CUDA-SingleNode.tar.gz](#) (Multi-copy GPU simulations. For multi-GPU machines use: "+devicesperreplica 1") (**March 24, 2022**)
- [NAMD_3.0alpha9_Linux-x86_64-multicore-CUDA-SingleNode.tar.gz](#) (Standard simulation.) (**February 28, 2021**)
- [NAMD_3.0alpha9_Linux-x86_64-netlrts-smp-CUDA-SingleNode.tar.gz](#) (Multi-copy GPU simulations. For multi-GPU machines use: "+devicesperreplica 1") (**February 28, 2021**)

https://www.ks.uiuc.edu/Research/namd/alpha/3.0alpha/download/NAMD_3.0alpha13_Linux-x86_64-multicore-CUDA-SingleNode.tar.gz

NAMD 3.0 Single-GPU-Per-Replicate Alpha Versions For Download:

The builds in this section support only a SINGLE-GPU offloading scheme. Builds for multi-GPU scaling per-replicate are posted in the section above.

NAMD 3.0 alpha 6 and later add support for GPU-accelerated FEP/TI simulations

- [NAMD_3.0alpha8_Linux-x86_64-multicore-CUDA-SingleNode.tar.gz](#) (Standard simulation.) (**December 21, 2020**)
- [NAMD_3.0alpha8_Linux-x86_64-netlrts-smp-CUDA-SingleNode.tar.gz](#) (Multi-copy GPU simulations. For multi-GPU machines use: "+devicesperreplica 1") (**December 21, 2020**)
- [NAMD_3.0alpha7_Linux-x86_64-multicore-CUDA-SingleNode.tar.gz](#) (Standard simulation.) (**October 16, 2020**)
- [NAMD_3.0alpha7_Linux-x86_64-netlrts-smp-CUDA-SingleNode.tar.gz](#) (Multi-copy GPU simulations. For multi-GPU machines use: "+devicesperreplica 1") (**October 16, 2020**)
- [NAMD_3.0alpha6_Linux-x86_64-multicore-CUDA-SingleNode.tar.gz](#) (Standard simulation.) (**August 12, 2020**)
- [NAMD_3.0alpha6_Linux-x86_64-netlrts-smp-CUDA-SingleNode.tar.gz](#) (Multi-copy GPU simulations. For multi-GPU machines use: "+devicesperreplica 1") (**August 12, 2020**)

NAMD 3.0 alpha 4 and later automatically manage stepsPerCycle and pairlistsPerCycle

- [NAMD_3.0alpha5_Linux-x86_64-multicore-CUDA-SingleNode.tar.gz](#) (Standard simulation.) (**July 22, 2020**)
- [NAMD_3.0alpha5_Linux-x86_64-netlrts-smp-CUDA-SingleNode.tar.gz](#) (Multi-copy GPU simulations. For multi-GPU machines use: "+devicesperreplica 1") (**July 22, 2020**)
- [NAMD_3.0alpha4_Linux-x86_64-multicore-CUDA-SingleNode.tar.gz](#) (Standard simulation.) (**July 18, 2020**)
- [NAMD_3.0alpha4_Linux-x86_64-netlrts-smp-CUDA-SingleNode.tar.gz](#) (Multi-copy GPU simulations. For multi-GPU machines use: "+devicesperreplica 1") (**July 18, 2020**)

NAMD 3.0 alpha 3 and earlier need stepsPerCycle and pairlistsPerCycle tuned, e.g. to 400 and 40 respectively, for best performance



文件



+ 代码

+ 义平

复制到云端硬盘

修改namd2为namd3



磁盘



{x}



GaMD_input_3HTB

NAMD_3.0alpha13_Linux-x86_64-m...

charmm-gui-6838781093

jz4

namd

restraints

toppar

README

step3_input.crd

step3_input.pdb

step3_input.psf

step4_equilibration.inp

step5_production.inp

sysinfo.dat

namd_gamd_total

toppar

3htb.cif

3htb_hetc.crd

3htb_hetc.pdb

3htb_proa.crd

3htb_proa.pdb

addCrystPdb.py

checkfft.py

<>



磁盘

可用存储空间: 143.37 GB

```
[5] charmm-gui-6838781093/toppar/par_all3
charmm-gui-6838781093/toppar/toppar_a
charmm-gui-6838781093/toppar/toppar_a
```

```
[8] 1 %cd /content/charmm-gui-68387810
/content/charmm-gui-6838781093/namd
```

Modify namd2 path in README, su
[/content/NAMD_Git-**-Linux-x86_](#)



1

Start a normal simulation (skip this)

```
[ ] 1 !csh README
```

Run GaMD

```
[ ] 1 %cd /content/charmm-gui-68387810
```

Substitute all "namd2" with namd3

README x

```
14 set equi_prefix = step4_equilibration
15 set prod_prefix = step5_production
16 set prod_step = step5
17
18 # Running equilibration step
19 /content/NAMD_3.0alpha13_Linux-x86_64-multicore-CUDA/namd3
20
21 # Running production for 10 nanoseconds
22 set cnt = 1
23 set cntmax = 10
24
25 while ( ${cnt} <= ${cntmax} )
26     # create appropriate input file using ${prod_prefix}.inp
27     if ( ${cnt} == 1 ) then
28         set outputname = "${prod_step}_${cnt}"
29         # change only the output name
30         sed "s/${prod_prefix}/${outputname}/" ${prod_prefix}
31     else
32         @ cntprev = ${cnt} - 1
33         set inputname = "${prod_step}_${cntprev}"
34         set outputname = "${prod_step}_${cnt}"
35         # change input and output names from template file
36         sed "s/${equi_prefix}/${inputname}/" ${prod_prefix}.
37         sed "s/${prod_prefix}/${outputname}/" > ${prod_
38     endif
39
40 # run the simulation for 1 nanosecond
41 /content/NAMD_3.0alpha13_Linux-x86_64-multicore-CUDA/n
42
43 @ cnt += 1
44 end
45
```

文件

🔍

📁

📄

👁

{x}

📁

..

namd_gamd_total

restraints

toppar

README

step3_input.crd

step3_input.pdb

step3_input.psf

step4_equilibration.inp

step5_production.inp

sysinfo.dat

toppar

3htb.cif

3htb_hetc.crd

3htb_hetc.pdb

3htb_proa.crd

3htb_proa.pdb

addCrystPdb.py

checkfft.py

checkfft.str

crystal_image.str

glycan.yml

input.config.dat

step1_pdbreader.crd

磁盘

可用存储空间: 143.36 GB

+ 代码

+ 文本

复制到云端硬盘

[] 1

Start a normal simulation (skip this)

🔄

1

!csh README

Run GaMD

[] 1

%cd /content/charmm-gui-68387810

Substitute all "namd2" with namd3

[] 1

!csh README

README

14

set prod_prefix = step5_production

15

set prod_step = step5

16

17

Running equilibration steps

18

/content/NAMD_3.0alpha13_Linux-x86_64-multicore-CUDA/namd3 \$

19

20

Running GaMD prerun for collecting parameters

21

set GaMD_prerun1 = step4.1_GaMD_init

22

set GaMD_prerun2 = step4.2_GaMD_upda

23

24

sed "s/\${prod_prefix}/\${GaMD_prerun1}/;s/#MDprep/100000/;s/#M

25

| s/#UPparm/0/;s/accelMDGRestart on/accelMDGRestart off/; \

26

| s/accelMDGRestartFile #Pstep.restart.gamd/" \${prod_pref

27

/content/NAMD_3.0alpha13_Linux-x86_64-multicore-CUDA/namd3 \$

28

29

sed "s/\${prod_prefix}/\${GaMD_prerun2}/;s/#MDprep/0/;s/#MDparm

30

| s/#UPparm/400000/;s/#Pstep/\${GaMD_prerun1}/" \${prod_pref

31

/content/NAMD_3.0alpha13_Linux-x86_64-multicore-CUDA/namd3 \$

32

33

sed "s/\${prod_prefix}/step5_GaMD_production/;s/#MDprep/0/;s/#

34

| s/#UPprep/0/;s/#UPparm/0/" \${prod_prefix}.inp > step5_GaM

35

set prod_prefix = step5_GaMD_production

36

37

38

Running production for 10 nanoseconds

39

set cnt = 1

40

set cntmax = 10

41

42

while (\${cnt} <= \${cntmax})

43

| # create appropriate input file using \${prod_prefix}.inp a

44

| if (\${cnt} == 1) then

45

| set outputname = "\${prod_step}_\${cnt}"

46

| # change only the output name